




An Introduction to Partitioning

Andrew Pruski

@dbafromthecold



Please silence
cell phones



Explore everything PASS has to offer

Free Online Resources
Newsletters
PASS.org



24HOURS
of  **PASS**

Free online webinar
events



PASS
LOCAL
GROUPS

Local user groups
around the world



 **PASS**
SQLSATURDAY

Free 1-day local
training events



PASS
VIRTUAL
GROUPS

Online special
interest user groups



 **PASS**
MARATHON

Business analytics
training



PASS
VOLUNTEERS

Get involved

Session evaluations

Your feedback is important and valuable.

Submit by 5pm Friday, November 16th to win prizes.

3 Ways to Access:



Go to passSummit.com



Download the GuideBook App
and search: PASS Summit 2018



Follow the QR code link displayed on session signage throughout the conference venue and in the program guide



Andrew Pruski

SQL Server DBA



/andrewpruski



@dbafromthecold

SQL Server DBA

Working with SQL Server for 10 years

Microsoft Data Platform MVP

Since December 2017

Based in Dublin, Ireland

Originally from Swansea, Wales. Living in Ireland for 5 years

Session Aim

To give you a base of knowledge to work with partitioning in SQL Server



Agenda

Partitioning Definition

Partitioning Key

Partition Functions & Schemes

Indexing Considerations

Splitting, Merging & Switching Partitions

Implementing Sliding Windows

Filegroup Restores



Partitioning Definition

Splitting a table horizontally into different units

Units can be spread across different physical locations

Limit of 15,000 partitions per table

Primarily for maintenance of data

Specialist functions available to manage data

Benefits

Partitioned tables appear as normal tables

Data is automatically mapped to the correct partition

Specialist operations allow for easy management of data

Individual partitions can be compressed

Individual partitions can be rebuilt

Drawbacks

- Requires management of partitions and filegroups
- Specialist operations can be blocked by DML operations
- Foreign keys referencing partitioned table will prevent switch operations
- Performance of queries not referencing the partitioning key will be affected

Creating a partitioned table



Partitioning key

Column in the table which defines partition boundaries

How is the data going to be split?

Archiving/retention policy for the data?

How is the table going to be queried?

All column types except timestamp, ntext, text, image, xml, varchar(max), nvarchar(max), or varbinary(max)

Partition Functions

Maps rows in the table to a partition

```
CREATE PARTITION FUNCTION [NAME] (DATATYPE)  
    AS RANGE RIGHT | LEFT  
FOR VALUES (n,n1,n2...nx);
```

Left / Right Range Types

Defines which side of the boundary the value specified belongs

```
CREATE PARTITION FUNCTION [MyPartitionFunction](DATE)
    AS RANGE RIGHT | LEFT
FOR VALUES (2016-01-01,2017-01-01,2018-01-01);
```

RIGHT

$2017-01-01 \leq x < \underline{2018-01-01}$

LEFT

$\underline{2017-01-01} < x \leq 2018-01-01$

Partition Schemes

Maps partitions to filegroups

```
CREATE PARTITION SCHEME [NAME]  
    AS PARTITION [FUNCTION NAME]  
[ALL] TO (FILEGROUP, FILEGROUP, FILEGROUP...);
```

Creating a partitioned table

```
CREATE TABLE dbo.PartitionedTable  
    (ID      INT IDENTITY(1,1),  
     ColA    VARCHAR(10),  
     ColB    VARCHAR(10),  
     CreatedDate DATE)  
ON PartitionScheme(CreatedDate);
```


Demo



Indexing



Clustered indexes

Create on the partition scheme specifying the partitioning key

Unique – the partitioning key has to be explicitly specified

Nonunique – the partitioning key will be added by SQL if not explicitly specified

Nonclustered indexes

An index that is created using the same partition scheme as the base table is *aligned*

An index that is created on a different filegroup or using a different partition scheme is *non-aligned*

Nonclustered indexes

Unique - the partitioning key has to be explicitly specified

Nonunique - the partitioning key will be added by SQL if not explicitly specified as an included column

Demo



Merging & Splitting



Merging Partitions

- Removes a partition

- Effectively “merges” two partitions into one

- Meta-data only operation if performed on an empty partition

- Data *will* be moved if partition is not empty, causing blocking and transaction log growth

Merging Partitions

```
ALTER PARTITION FUNCTION [NAME]()  
    MERGE RANGE (VALUE);
```

Splitting partitions

- Creates a new partition with new boundary value

- New boundary value must be distinct from other values

- Takes a schema modification lock on the table

- Meta-data only operation if partition is empty

- SQL *will* move data to the new partition if the data crosses the new boundary value

Splitting partitions

```
ALTER PARTITION SCHEME [NAME]  
    NEXT USED [FILEGROUP];
```

```
ALTER PARTITION FUNCTION [NAME]()  
    SPLIT RANGE (VALUE);
```

Demo



Switching



Switching partitions

Move a partition from one table to another

Meta-data operation, runs immediately

Both tables must have the same structures

Destination partition must be empty or...

if destination table is not partitioned, it must be completely empty

Switching partitions

```
ALTER TABLE [Source Table]  
    SWITCH PARTITION Partition_Number  
TO [Destination Table]  
    PARTITION Partition_Number;
```


Demo



Sliding Windows



Partition Sliding Windows

Method to remove old data and bring in new data periodically

Implements the SWITCH, MERGE, & SPLIT functions

Partitions in the table move “forward” but the overall number of partitions remains the same

Partition Sliding Windows

1. SWITCH oldest partition in live table to archive table
2. MERGE oldest partition
3. SPLIT new partition
4. Load new data into staging table
5. SWITCH data from staging table to live table
6. Update statistics on live table

Demo



Filegroup Restores



Filegroup Restores

Can be useful for VLDBs

Can be used to restore live partitions to development

Individual partitions are on different filegroups

Data in older partitions does not change or is not needed

Reduce recovery time for “active” data

Demo



A quick story



Questions?





Thank You

Learn more from Andrew Pruski



@dbafromthecold



dbafromthecold@gmail.com

